# sHeap: Mitigating Dangling Pointers

Matt McCreesh
*Stevens Institute of Technology*
mmccrees@stevens.edu

Thomas Pyle
*Stevens Institute of Technology*
tpyle@stevens.edu

Connor Zapfel
*Stevens Institute of Technology*
czapfel@stevens.edu

*Abstract*—When an object in memory is deallocated, the freed memory can still be accessible by the program if any pointers to it are not nullified. Pointers left pointing to objects that have been freed are referred to as dangling pointers. They point to a location that once had a valid object of the type they point to but no longer does. The presence of such pointers can lead to use-after-free vulnerabilities. These vulnerabilities can be just as worrisome as spatial memory errors like buffer overflows, as they can be used by an attacker to hijack the control flow of an application. In this paper, we propose and evaluate a runtime allocator, sHeap, which mitigates the dangers of use-after-free vulnerabilities by preventing arbitrary code execution through dangling pointers. Our solution is based off ideas and approaches proposed and implemented by Cling in [1].

*Index Terms*—Memory allocator, use-after-free, heap, memory safety

## I. INTRODUCTION

Dangling pointers to the heap, pointers left pointing to deallocated objects, are a major security concern. They were once thought to be too difficult for attackers to exploit, as attackers would need to be able to force a program to put his data in the same location as the dangling pointer pointed to. Because of this, in the past vulnerabilities were sometimes left unpatched for long periods of time and written off as non-critical. But as spatial memory protections became more popular, the arms race between attackers and defenders continued and it was shown that attackers could reliably exploit use-after-free vulnerabilities. Heap Spraying and Heap Feng Shui allows an attacker to reliably align his own data with that of a dangling pointer. This means that reliable exploitation of these vulnerabilities is possible. In fact, these vulnerabilities have been showing up in the wild, including in a zero-day exploits on the Google Chrome web browser, where a vulnerability was patched in March 2019. Use-after-free vulnerabilities are very attractive for attackers to exploit because there is potential for control flow hijacking. If a dangling pointer to a heap object with function pointers exists, an attacker may be able to launch code-reuse attacks, such as return-to-libc, by convincing the program to make a function call through a dangling pointer.

An interesting target of attacks is C++ objects. The C++ language is object based and allows for inheritance. In C++, a parent class can declare some functions as virtual functions, which can be redefined in derived classes. Virtual functions should execute the version of the function of the object that is allocated, as opposed to the type of variable pointing to it. In order to do this so called virtual function dispatch, C++ objects with virtual functions start with a pointer (called the vptr) to a virtual table (vtable), which stores function pointers for virtual functions of this object. When a virtual function of a C++ object is called, the vptr is used to find the vtable, which is used to find the function pointer of the correct function. Use-after-vulnerabilities can exploit this virtual function dispatch mechanism. If an attacker controls the buffer in the heap that overlaps with memory that once contained a C++ object, and there is a dangling pointer to that object, the attacker can overwrite what was the vptr field was and set it to point to his own counterfeit vtable. Then, when the dangling pointer tries to execute a virtual function, it will instead execute whatever the attacker points his vtable to.

Unfortunately, use-after-free vulnerabilities have proven very difficult to defend against. They are hard to detect with static analysis and code reviews, as they require understanding the state of memory when analyzing pointer dereferences. Such memory state depends on code that has already executed, and synchronization of code execution, making this task especially difficult in large applications. To address this problem, we introduce and implemented a memory allocator, sHeap. Memory allocator changes have been made to address other security issues such as double frees and overwrites of heap metadata, so there is precedent in implementing security hardening in the run time allocator itself. Our solution is based off of ideas proposed in [1]. sHeap takes many elements from Cling, though it does not implement every optimization of Cling, as some of the specific optimizations for small allocations are left for future work.

sHeap redefines the memory allocator interfaces (malloc(), new(), calloc(), and realloc()) to allow for pooling of allocations over time with objects of the same type. sHeap uses the call site to malloc, calloc, and realloc to determine the type, assuming that a call site to malloc or new will always request memory of the same type. It also does malloc wrapper detection to find call sites to simple malloc wrappers so as to avoid pooling objects that are of different types because they are allocated from wrappers of malloc. Simple malloc wrappers may just check the value returned by malloc or collect statistics before returning what malloc returns. They obscure the real allocation site, sHeap uses to determine allocation type, so they must be unwrapped. This wrapper detection allows us to find the call site to the new C++ operator, which allows us to properly pool C++ objects and dynamically allocated arrays.

Our allocator defends against arbitrary code execution through a dangling pointer. It does not prevent dangling

pointers from being used, but it does ensure that they always point to objects of the same type. Therefore, if an attacker tries to execute a function call through a dangling pointer, the code executed should be a valid function for an object of that type. sHeap is not concerned with data only attacks where fields are overwritten.

Ultimately, our solution is effective from a security perspective, but it has high CPU and memory overhead for some of the benchmarks we used. This is likely a result of the fact that sHeap is not currently optimized for small allocations. We focused on supporting allocations of any size to make it work for as many programs as possible, so we tailored our project to support large allocations instead of also being optimized for small allocations. The use of buckets for small allocations should allow for great performance benefits for applications making many small allocations. This will result in less system calls and also significantly reduce fragmentation. We leave such performance improvements for future work.

## II. BACKGROUND

### A. Current Heap Implementations

The most common heap implementation currently in use is GNU malloc, also known as dlmalloc. GNU malloc works by allocating blocks of data using mmap calls, as well as brk calls, in order to find unavailable blocks to give to the use. Malloc does this through the use of a free list which keeps track of the size of program-owned memory that is no longer attached to allocated blocks. Most heap implementations, including GNU malloc, do not make any effort to mitigate use-after-free vulnerabilities.

### B. Related Work

A naive approach to defend against dangling pointers is by preventing any heap allocated objects from ever sharing the same memory space. This is done by never reusing freed memory, effectively replacing free() with a function which does nothing. This clearly will have high physical memory overhead for long running applications. More intelligent implementations of the naive solution aim to signal the operating system to release physical pages while not reusing virtual addresses, but this still can still have high physical memory overhead because of fragmentation, where small allocations hold back pages from being released. The naive solution is simple implement, but incurs an extremely large overhead of consumed virtual memory.

Some other defenses include probabilistic defenses, such as DieHard [2] and Archipelago [3]. These defenses both rely on defeating use-after-free attacks by making it improbable that the attacker will be able to align heap objects. It does this by randomly locating objects in a larger heap than usual. This of course results in worse memory consumption, because the heap needs to be larger in order to ensure the probabilistic defense.

Cling [1] is a memory allocator which attempts to provide protection from dangling pointer in particular through the use of type safety in heap allocations. It does this by analyzing the return site of malloc (by looking at malloc's stack frame), as well as categorizing objects based off this size. It does this in order to determine run-time type information based off assumptions about the way that programmers frequently perform allocations of data. Cling performs many optimizations and as such is capable of reducing runtime and memory overhead to something comparable (and even sometimes better) than traditional GNU malloc.

## III. DESIGN

### A. Objective

sHeap is designed to protect heap memory in a manner similar to Cling, but unfortunately imposes a number of limitations due to the time and resource constraints of the team. This means that although Cling is capable of optimizing small allocations by storing many in a single allocated block [1], we were unable to perform this optimization in order to prioritize goals such as the detection of wrappers and thread safety.

The implementation of this technique was done by replacing the 4 malloc calls defined in traditional glibc. These can all be performed by loading our shared library before loading stdlib. This is done by using the environment variable LD_PRELOAD. By redefining these 4 methods, we ensure that any programs which use our heap allocator will always have a consistent view of the heap through our allocator.
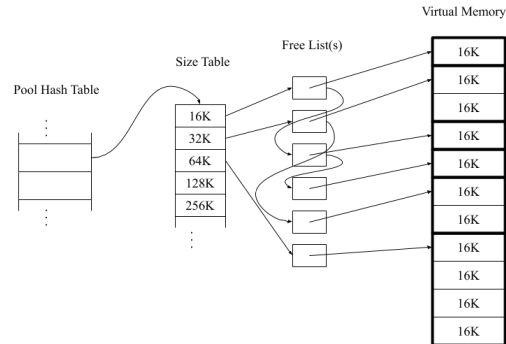
### B. Architecture



Fig. 1: sHeap Architecture Summary.

sHeap is primarily made up of three coordinating components. First, the pool hash table tracks the mapping from allocation sites to their corresponding site pools. Then the site pool manages the different pointers to the free space available for each size class. The free space is stored in a linked list structure called a free list. 1 provides a visual representation of the sHeap architecture. In the next three subsections, we will go into the design and purpose of the three main components mentioned above.

*1) Pool Hash Table:* The first layer of interaction into the sHeap ecosystem is the pool hash table (PHT). It has a single purpose: to map allocation sites to their appropriate site pool. A hash table allows for constant time average case access complexity which is essential for maintaining low overhead. The PHT makes use of the allocation site in the hash function making for an efficient and collision-resistant hash function. The team opted for linear probing as the collision handling strategy for simplicity of implementation, and it has proven effective in practice. Each entry in the hash table at a minimum needs to store a pointer to the isolated allocation site pool in order to connect the allocation site to the next stage in the pipeline.

*2) Site Pool:* The next primary data structure in sHeap is the site pool, also referred to as the size table. This is a struct for every allocation site, managing the metadata for each allocation site pool. The primary job of this data structure is to manage the pooling of allocations by size class, which is why it is referred to as the size table. A size class is a power of 2 KB allocation size starting at 16KB. That is, every allocation size of 16KB or less will be in the 0th size class and be allocated only 1 block. Any allocation size of 16KB to 32KB will be in the first size class and be given 2 blocks. The next size class corresponds to an allocation of 4 blocks and is for allocation sizes between 32KB and 64KB. The size table structure stores a table of pointers to free lists heads for each size class, as well as metadata about allocation sizes used to determine if a call site is a wrapper. Every time a new allocation site is found, a site pool element is created for that allocation site. For future allocation requests, this structure will be used to find the free list head for the size class being requested so type safe address space reuse can be done. During future deallocations, this table is used to append freed block nodes to the correct free list for the allocation site and size class of the original allocation. We will get more into the specifics of this when discussing the free lists themselves. The site pool structure is also very useful for wrapper detection. By having a field that stores the original size requested, as well as whether or not this allocation site was determined to be a wrapper or a non-wrapper, we can correctly and efficiently find the real allocation site of malloc wrappers, and know when to suspect an existing allocation site is a wrapper. We will discuss the importance of this in the next section.

*3) Freelists:* The free list is designed to keep track of previously allocated, but freed, blocks. In order to preserve the mitigation described, it is necessary for the free list to be dependent on the allocation site and the size class of allocations, and as such each pair of allocation and size class has its own free list. The head of the free list is stored in the size table, and allocation attempts to find a previously allocated node before it attempts to request more allocated memory. The free list is designed to make reverse lookups - programmed owned data to free list nodes - performed in constant time. This is so that the free operation is performed in constant time, and is capable of performing a lookup for the relevant size class and type information stored in the node.

This is necessary in order to append the newly freed node to the correct free list.

## IV. Implementation

### A. Pool Hash Table

As discussed in section 3, the pool hash table exists to store the mapping of allocation sites to their site pools. This requires the implementation of a hash table that is stored in the out-of-band metaheap. We implemented the hash table as a single contiguous array of pht_entry structs. Each of these structures stores exactly two pieces of information: the allocation site and the pointer to the corresponding site pool. We have to store the allocation site in the pht_entry in order to detect collisions. If two allocation sites end up hashing to the same hash index, there is no way to confirm correctness unless we store that allocation site. Our PHT was not dynamically scalable which means the size of the PHT is decided at initialization by a constant and doesnt change throughout the runtime of sHeap. This is something that could be improved in the future in order to accomodate more allocation sites, but in our testing, this was not an issue that we encountered. The hashing algorithm that was used is pretty simple. Since the input is an eight-byte pointer (the allocation site), we just isolate xor the upper and lower four bytes then mod the result by the size of the PHT as dictated by that constant variable from earlier.

### B. Site Pool

Once the PHT has found the site pool for this allocation, which stores the pool metadata for this allocation site, we can discuss how allocation works. On the first request from an allocation site, the allocation site will be inserted into the PHT but the size table pointer for the pool metadata will be null. This will trigger the size table to create a new element for the allocation site and update the PHT element to have a pointer to that struct. In future allocations from this allocation site, the size table will try to utilize the free list for the given size class by passing in the location that it stores the free list head pointer to the allocation function. That function, which is discussed further in the next section, will use the freelist for allocations if it is not empty. The important point is the site pool size table allows sHeap to efficiently find the beginning of the correct free list so blocks can be reused if the free list is not empty.

### C. Freelists

The free list is implemented as a structure containing all of the relevant type information of a block of memory, as well as pointers to the next and previous nodes in the free list. The actual allocation operation is passed through the pool hash table and the size table, and then arrives at the free list in order to determine whether or not to create a new allocated block, or whether to reuse a previously freed block. This is done by checking whether the flist passed by the size table contains a valid pointer (in which case there is a block to use), and if not, it will push the program break and claim a new one. When performing the free operation, sHeap needs to

lookup the relevant information for the node (its allocation site and its size class) in order to find the correct free list to use. It does this by preserving the order of the program data with the free list nodes, so the offset of the program data from its start (in blocks) corresponds to the free list node to use. This allows constant-time reverse lookups of the free list nodes. Once we perform the reverse lookup, we are able to use the information stored in the node to append it to the appropriate free list, by looking up the allocation site in the pool hash table, and then the corresponding size in the size table. The node is then appended to the front of that free list, to be reused in a subsequent malloc call.

### D. Wrapper Detection

At this point we have shown how our memory allocator works. The security benefits we achieve all stem from constraining memory reuse to be within pools of the same type, and aligning those allocations to start at the same positions. Unfortunately, using the call site to malloc as the allocation site is not always sufficient for uniquely identifying the type of the allocation. Simple malloc wrappers can exist in programs and this would subvert our protections. A malloc wrapper is a function that calls malloc on behalf of a caller, simply returning the pointer it mallocs. We found that the new C++ operator was a malloc wrapper that we needed to be able to handle for type safe memory reuse for C++ programs. Other common uses of malloc wrappers include checking if the return is null before sending it to the caller and collecting statistics on allocations. In order to accurately pool allocations by type, we must find the actual allocation site. This is the call site to a simple wrapper if one exists. sHeap uses allocation sizes to determine whether it suspects and allocation site is really a wrapper or not. On the first allocation, it stores the requested allocation size in the size table for the allocation site pool. On future allocations from that call site, if the size requested is different, sHeap suspects it may be a simple allocation wrapper. In order to determine if a call site is a wrapper, sHeap must find out if the function calling it returns the pointer returned by malloc. In order to do this, when sHeap suspects an allocation site is a wrapper, it uses libunwind to find the return address on the stack frame of the caller to malloc, and it overwrites that return address with an assembly routine that will mark the allocation site as a wrapper or a non-wrapper. Malloc then does allocation like normal, saving values that will be used for wrapper detection to global variables. These values include the return value that was overwritten on the stack, the location that was overwritten on the stack, and a pointer to the allocation site size table so the call site to malloc can be marked as a wrapper or non-wrapper. Malloc returns to the suspected wrapper, which will eventually return to the assembly routine. The assembly routine operates in the stack frame of the caller to the suspected wrapper, so it is very important not to overwrite memory in the stack frame. The routine checks the value of %rax (the return value of the suspected wrapper) and compares it to the last malloc return. If they are equal, it uses the archived size table pointer to

mark the call site as a wrapper, and if they are not equal it marks that call site as a non-wrapper. It then restores the value of %rax and all callee saved registers and jumps back to the return address the wrapper normally would have returned to, allowing the program to go back to its normal flow. Once a call site to malloc is marked as a wrapper, every time that call site calls malloc, malloc will lookup that it is a wrapper and use libunwind to unwind the stack 2 layers to get the actual allocation site, the call site to the wrapper.

One challenge with wrapper detection is when there are multiple calls to malloc before the suspected wrapper returns. This can be the result of a suspected wrapper function calling malloc multiple times, or a suspected wrapper function calling another function that calls malloc. In either case, our simple malloc detection mechanisms rely on global variables to store values returned by malloc, return addresses, and locations to mark as a wrapper or non-wrapper. When doing wrapper detection this way, we can have a situation where two call sites become suspected as wrappers before the first suspected wrapper returns. To handle this, we abort detection of the first suspected wrapper when a second suspected malloc wrapper arises before the first suspected wrapper returns. We believe this is a valid approach, as we are trying to identify simple wrappers and it is unlikely that a simple wrapper would make multiple calls to malloc. A similar approach is taken by [1], which assumes that a function making multiple malloc calls is not a malloc wrapper.

### E. Thread Safety

Our last major challenge was making our allocator thread safe. Because we had several data structures that needed to be made thread safe, including free lists and table accesses, this was not a trivial problem. Another added challenge was the fact that our data and metadata had to maintain alignment. This required synchronization of user-heap and meta-heap allocations. When a new block was acquired from sbrk or a free list, we had to associate the correct metadata node with that block. In order to do this safely without introducing concurrency errors, we took a fairly naive approach to thread safety. We used libpthread mutexes to protect entire function calls, so that the entire allocation routine was considered the critical region and only one thread could execute it at a time. This clearly would have a performance penalty for programs with many threads that make allocation calls. A future direction would be to minimize our critical region protected by mutexes and to formally model and verify the thread safety of our future implementation. Our current version of sHeap does guarantee thread safety, in terms of preventing race conditions while not introducing deadlock, but it does so with less than ideal performance. Making our allocations and deallocations thread safe was one challenge, but a more interesting problem was ensuring thread safety for our wrapper detection mechanisms. In order to do wrapper detection and to mark call sites as being wrappers, certain values needed to be saved in global variables. We needed a way to prevent threads from overwriting the wrapper detection variables used by

other threads. Using semaphores to do this was unreasonable because the critical region would be exited only after the suspected wrapper returned. One approach we considered was using registers to propagate the values we were storing in global variables, but this is not reliable as overwriting a callee saved register can break some programs, and overwriting a caller saved register may not propagate to the assembly routine. So in order to guarantee that we could reliably detect wrappers, we made our wrapper detection global variables be thread local variables, so each thread operated on its own values without overwriting the value at other threads. With this approach, we can reliably and efficiently detect wrappers, even if the program has multiple threads having suspected wrappers simultaneously.

## V. Evaluation

All evaluation took place on a 2014 HP ZBook 15 running Ubuntu 18.04 LTS with an Intel i7-4700MQ clocked at 2.4GHz and equipped with 8GB of RAM. We tested four programs for both run time and memory overhead and two programs for just memory overhead. Python 3, ping, espresso, and GNU Make belong to the former, while evince and Gimp belong to the latter. With Python 3, a small script was constructed that simply bulk allocated simple Point objects and then appended them to a list. In total we allocated nine million of these. While the team has no reason to suspect that ping will give any incredible benchmarking results as an I/O-bound program, it was included as a sanity-check for evaluation purposes. We simply pingd google.com five times. Espresso is a heuristic logic minimization tool from [4] that was used in the testing of Cling due to its allocation patterns so we felt it best to include it here. GNU Make is a real-world program that sees active usage. For evaluation, Make was tested by building Espresso from source. Gimp and evince were included in the benchmarking process because we wanted to perform some testing of bigger allocations ($> 16KB$). Evince is a PDF reader that ships with Ubuntu 18, and Gimp is an image manipulation program that, unsurprisingly, operates on images. Both images and PDFs tend to be bigger in size. Both Gimp and Evince were tested by simply opening an image or PDF to measure the memory overhead for these larger allocations. It is important to note that all benchmarks were tested by making use a shell script in order to properly record the metrics discussed below. This means the raw numbers will include any memory overhead that was also taken on by the script, but both with and without sHeap have this same taint.

### A. Memory Overhead

We captured two metrics when looking at memory overhead: resident set size (RSS) and virtual memory size (VSize). Both RSS and VSize were recorded with respect to their peak values because this is what the Cling authors did, and it seems reasonable for benchmarking purposes. We captured this overhead using a Python tool called Syrupy.py ( [5]) that queries /proc/[pid]/status at a specified interval and outputs

both the RSS and VSize which was then recorded by the team. 2 shows the results of the raw VSize measurements.



Fig. 2: sHeap Virtual Memory Overhead.

sHeap has relatively minimal overhead in the VSize measurement for Gimp, evince, ping, and even Python 3. Espresso and Make have more of a discrepancy in this area, and the normalized results will be seen at the end of this subsection. RSS is a slightly similar story, as again both Make and Espresso appear worse, but to a greater intensity this time.



Fig. 3: sHeap Resident Set Size Memory Overhead.

The other four programs all appear to be in a similar range of overhead for RSS. If we now look at 3, we can see how much memory overhead is being incurred with sHeap now.

Both Make and Espresso have significant RSS overhead, 2,691% and 4,069% respectively. Python 3 however sits more comfortably at around 12% RSS, while the remaining three programs are all under 1%. With respect to VSize, Espresso is the only poorly performing benchmark coming in at 324%.

### B. Runtime Overhead

Now we turn out attention to the run-time overhead that is brought upon by the usage of sHeap. There are only four programs at work here as described in the section 5 introduction. Again, there are just two metrics that were tracked in the runtime department: User time and System time. If we look at

Fig. 4: sHeap Memory Overhead.


Fig. 6: sHeap User Time Overhead.

the raw user time performance between sHeap and non-sHeap executions, we see sHeap performing worse across the board. This is especially true in the case of Make and Espresso, and less so in the case of Python 3 and ping. Ping logically makes sense as it is primarily bound by the I/O delays inherent to the nature of ping. System time is much worse in the case of Espresso and Make, where Espresso resulted in 4,913.33% and Make yielded 515.79%. This is most likely due to the lack of optimization for smaller allocations which forces a significantly higher rate of system calls to occur. Ping has a 100% increase in system time due to the fact that it makes a single malloc call, and sHeap has to make one on initialization of its metadata which is doubling the number of allocation calls to the kernel.


Fig. 5: sHeap System Time Overhead.

### C. Performance Hurdles

The team understands where sHeap is having its performance hindered most severely. Due to the time constraint on the project, we had to choose which features/capabilities to tackle and give priority. In a perfect world, small allocations ($< 16KB$) are performed in a more optimal manner. For example, the allocation of a small 16-byte structs should all be tucked together into a single virtual memory block, but as it stands now, a new block is allocated for each small allocation.

This means both a syscall for each small allocation as well as the wasted virtual memory from only using 16 bytes from a whole 16KB block. Implementing a more efficient solution to smaller allocations would greatly increase performance from both the time and memory overhead points of view. Another side effect of this that relates primarily to time benchmarking is the lack of quality caching. Because allocations (especially small ones) are isolated into their own memory blocks, the principle of locality that drives the benefits of using a cache is all but useless. Fragmentation with sHeap is high, which means caches arent able to effectively perform their task.

## VI. LIMITATIONS AND ISSUES

Our implementation introduces a number of limitations in order to ensure its functionality, as well as due to its prototype state. The most obvious limitation is that our heap manager relies on virtual memory being contiguous, and as a such all programs which use it cannot call sbrk in between malloc calls. This is because any such calls to sbrk will disturb the order, so the free list nodes don't match up with the appropriate allocated block. Programs that do use sbrk are detected however, and will cause the program to exit. Additionally, this system is also only applicable to 64 bit systems, due to the inefficiency of memory consumption, as well as due to the hash algorithm that was used in our pool hash table which expects 64 bit addresses. As mentioned in section 5-c, the implementation of improved handling for allocations less than 16KB is a big problem from a performance standpoint. This should be addressed in the future to make sHeap more usable. Another limitation is that sHeap requires a single large upfront allocation to store the out-of-band metaheap information and structures. This introduces a problem because now there is a restriction on how big the metaheap can grow (because of the continuity of memory issue from earlier). Further, in order to improve the performance of malloc calls, wrappers are only detected on subsequent calls, so the initial wrapper call is not detected as unusual, and as such it will cause the first one to have an allocation site in the wrapper. Further, we also can only detect the presence of a single layer of wrappers, and multiple layered wrappers are ignored - this is primarily due

to the fact that we consider it exceptionally unusual to use multiple layers of malloc wrappers.

## VII. Conclusion

Dangling pointers represent a significant vulnerability space for attacks in the modern age. In this paper we have presented sHeap, a memory allocator designed to safely reuse heap space in order to prevent an attacker from being able to align heap objects in order to arbitrarily execute code through the use of dangling pointer exploits. It is able to do this by reliably identifying objects accurately based off of their allocation site, in a manner that is thread safe.

## References

[1] Akritidis, P. "Cling: A Memory Allocator to Mitigate Dangling Pointers," *19th Usenix Security Symposium*, Washington DC, 2010. [Online]. Available: https://www.usenix.org/legacy/event/sec10/tech/full_papers/Akritidis.pdf.

[2] Berger, E., and Zorn, B. "DieHard: Probabilistic Memory Safety for Unsafe Languages," *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation*, ACM, NY, 2006, pp. 158–168. [Online]. Available: https://people.cs.umass.edu/ emery/pubs/fp014-berger.pdf.

[3] Lvin, V., Novark, G., Berger, E., and Zorn, B. "Archipelago: Trading Address Space for Reliability and Security," *SIGOPS Oper. Syst.*, Rev. 42, 2008, pp. 115–124. [Online]. Available: https://people.cs.umass.edu/ emery/pubs/asplos147-lvin.pdf.

[4] galengold, Espresso Logic Minimizer, 2017, GitHub Repository: https://github.com/galengold/espresso-logic.

[5] jeetsukumaran, Srupy: System Resource Usage Profiler, 2008, GitHub Repository: https://github.com/jeetsukumaran/Syrupy.